

# CHALMERS



## Rymd

Distributed encrypted peer-to-peer storage

*Bachelor's thesis in Computer Science*

NIKLAS ANDRÉASSON

ROBIN ANDERSSON

JOHANNES RINGMARK

JOHAN BROOK

ROBERT EDSTRÖM

Department of Computer Science and Engineering

*Division of Software Engineering*

CHALMERS UNIVERSITY OF TECHNOLOGY

Göteborg, Sweden 2014

Bachelor's thesis 2014:01



BACHELOR'S THESIS IN COMPUTER SCIENCE

# Rymd

Distributed encrypted peer-to-peer storage

NIKLAS ANDRÉASSON  
ROBIN ANDERSSON  
JOHANNES RINGMARK  
JOHAN BROOK  
ROBERT EDSTRÖM

Department of Computer Science and Engineering  
*Division of Software Engineering*  
CHALMERS UNIVERSITY OF TECHNOLOGY  
Göteborg, Sweden 2014

Rymd  
Distributed encrypted peer-to-peer storage  
NIKLAS ANDRÉASSON  
ROBIN ANDERSSON  
JOHANNES RINGMARK  
JOHAN BROOK  
ROBERT EDSTRÖM

© NIKLAS ANDRÉASSON , ROBIN ANDERSSON , JOHANNES RINGMARK , JOHAN  
BROOK , ROBERT EDSTRÖM, 2014

Bachelor's thesis 2014:01  
ISSN 1654-4676  
Department of Computer Science and Engineering  
Division of Software Engineering  
Chalmers University of Technology  
SE-412 96 Göteborg  
Sweden  
Telephone: +46 (0)31-772 1000

Chalmers Reproservice  
Göteborg, Sweden 2014

## ABSTRACT

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

Keywords: Some stuff, More stuff, Stuff

## SAMMANFATTNING

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Ut purus elit, vestibulum ut, placerat ac, adipiscing vitae, felis. Curabitur dictum gravida mauris. Nam arcu libero, nonummy eget, consectetur id, vulputate a, magna. Donec vehicula augue eu neque. Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Mauris ut leo. Cras viverra metus rhoncus sem. Nulla et lectus vestibulum urna fringilla ultrices. Phasellus eu tellus sit amet tortor gravida placerat. Integer sapien est, iaculis in, pretium quis, viverra ac, nunc. Praesent eget sem vel leo ultrices bibendum. Aenean faucibus. Morbi dolor nulla, malesuada eu, pulvinar at, mollis ac, nulla. Curabitur auctor semper nulla. Donec varius orci eget risus. Duis nibh mi, congue eu, accumsan eleifend, sagittis quis, diam. Duis eget orci sit amet orci dignissim rutrum.

Nam dui ligula, fringilla a, euismod sodales, sollicitudin vel, wisi. Morbi auctor lorem non justo. Nam lacus libero, pretium at, lobortis vitae, ultricies et, tellus. Donec aliquet, tortor sed accumsan bibendum, erat ligula aliquet magna, vitae ornare odio metus a mi. Morbi ac orci et nisl hendrerit mollis. Suspendisse ut massa. Cras nec ante. Pellentesque a nulla. Cum sociis natoque penatibus et magnis dis parturient montes, nascetur ridiculus mus. Aliquam tincidunt urna. Nulla ullamcorper vestibulum turpis. Pellentesque cursus luctus mauris.

Nulla malesuada porttitor diam. Donec felis erat, congue non, volutpat at, tincidunt tristique, libero. Vivamus viverra fermentum felis. Donec nonummy pellentesque ante. Phasellus adipiscing semper elit. Proin fermentum massa ac quam. Sed diam turpis, molestie vitae, placerat a, molestie nec, leo. Maecenas lacinia. Nam ipsum ligula, eleifend at, accumsan nec, suscipit a, ipsum. Morbi blandit ligula feugiat magna. Nunc eleifend consequat lorem. Sed lacinia nulla vitae enim. Pellentesque tincidunt purus vel magna. Integer non enim. Praesent euismod nunc eu purus. Donec bibendum quam in tellus. Nullam cursus pulvinar lectus. Donec et mi. Nam vulputate metus eu enim. Vestibulum pellentesque felis eu massa.

# TERMINOLOGY

## General terminology

**Adobe PhoneGap** A software enabling development of cross-platform hybrid smartphone applications - applications developed using web technologies, but packaged as native smartphone binaries.

**API** Application Programming Interface. An interface that software developers can use to get easy access to data and/or particular functionality for their software. Typically exposed as a software library or HTTP service.

**Bitcoin** The first and biggest widespread cryptocurrency.

**CA** Certificate Authority. A third party that is trusted to verify the validity of public keys and certificates.

**Chrome Apps** Similar to Adobe PhoneGap, but for desktop applications running in a Google Chrome sandbox.

**Cloud storage** A service that hosts data externally with seamless access over the internet.

**Cryptocurrency** A network transaction system that uses a fully distributed cryptographically secured ledger (called *blockchain*) to track transactions. The vast majority of cryptocurrencies are forks off Bitcoin.

**Centralized system** A system which has several nodes connecting to and depending on one or a few central endpoints.

**Decentralized system** A system where responsibilities are shared across the nodes and does not depend on a single, central endpoint.

**DHT** Distributed Hash Table. A notion in computer science of a distributed key-value store.

**Firefox OS** A smartphone operating system where all applications are web-based.

**GUID** Globally Unique Identifier. Used as pseudo-unique identifiers, such as keys in a database. Usually 128-bit values stored as 32 hexadecimal in groups separated by hyphens.

**NoSQL** All database systems which are not modelled in tabular relations. Examples are graphs, trees, and key-value stores.

**OpenPGP** A standard for data encryption and signing, originally coming from the proprietary software Pretty Good Privacy (PGP) and widely spread through the free implementation GPG.

**P2P** Peer-to-peer. Distributed, direct communication.

**PKI** Public Key Infrastructure. A system that associates (unique) user identities with their public keys. Typically implemented as a Web of Trust or with one or several CAs. Typically, trusted parties use their private keys to sign the public keys of users to verify the connection between an identity and a public key.

**RSA** A widely used public-key cryptosystem. As such, it builds on pairs of private and public keys where encryption with one can be reversed by decrypting with the other. This system is asymmetric and the security relies on the practical difficulty of factoring the product of two large prime numbers.

**RDBMS** Relational Database Management System, a popular type of database management system based on the relational model.

**SQL** Structured Query Language. A language for managing, querying and manipulating data in relational database systems.

**SQL injection** A technique for injecting malicious SQL code into the executing database queries in order to, for instance, dump the contents of the database.

**XSS** Cross Site Scripting is the technique for injecting client side scripts into web pages.

**Web of Trust** A type of distributed PKI that builds on peer-to-peer trust. The idea is that if Alice trusts Bob, then Bob is trusted introduce new identities and public keys for Alice.

## Terms with specific meaning in the Rynd project

**Identity** A unique, memorable string identifying a user within the network.

**Module** A delimited area of interest and functionality in system architecture.

**Node** A client in the network (such as a web browser).

**Resource** A file or folder in the network (a thing that can be shared between nodes)

**Rynd** The developer library for web based peer-to-peer sharing – the main product. Is also the Swedish word for *space*, which encompass the main ideas of the project.

**Shuttle** A web based file sharing application implemented using Rynd to demonstrate the basic capabilities of the project.



# Contents

<b>Abstract</b>	<b>i</b>
<b>Sammanfattning</b>	<b>ii</b>
<b>Terminology</b>	<b>iii</b>
<b>1 Introduction</b>	<b>1</b>
1.1 Background . . . . .	1
1.2 Purpose . . . . .	2
1.3 Problem . . . . .	2
1.3.1 Decentralization of system logic . . . . .	3
1.3.2 Source code integrity . . . . .	3
1.3.3 Peer identity verification . . . . .	3
1.3.4 Resource storage . . . . .	3
1.3.5 Resource identification . . . . .	3
1.3.6 Communication flow and transfer initiation . . . . .	4
1.4 Scope . . . . .	4
1.5 Similar technologies . . . . .	4
<b>2 Methodology</b>	<b>5</b>
2.1 Evaluation of technologies . . . . .	5
2.2 Prototyping . . . . .	6
2.3 Implementation . . . . .	6
<b>3 System design</b>	<b>6</b>
3.1 Overall system design . . . . .	6
3.1.1 Resources . . . . .	6
3.2 Peer identity verification . . . . .	6
3.3 Decentralization . . . . .	7
3.4 Modularity . . . . .	7
<b>4 System implementation</b>	<b>8</b>
4.1 Data storage . . . . .	8
4.1.1 IndexedDB . . . . .	9
4.1.2 Implementation . . . . .	11
4.2 Communication . . . . .	12
4.2.1 WebRTC . . . . .	12
4.2.2 Implementation . . . . .	13
4.3 Authentication . . . . .	15
4.3.1 Namecoin . . . . .	15
4.4 Encryption . . . . .	15
4.5 Testing . . . . .	15

<b>5</b>	<b>Results</b>	<b>15</b>
5.1	Source code . . . . .	16
<b>6</b>	<b>Discussion</b>	<b>16</b>
6.1	Quality of the resulting product . . . . .	16
<b>7</b>	<b>Conclusion</b>	<b>16</b>
	<b>References</b>	<b>17</b>
	<b>Appendices</b>	<b>19</b>
<b>A</b>	<b>Communication flow</b>	<b>I</b>
<b>B</b>	<b>List of external libraries</b>	<b>I</b>

# 1 | Introduction

IN A DAY and age where people all over the world own more than one digital device, there is a growing need for services which let users easily store, access, sync and share personal files. Due to the Internet it is possible to store data *in the cloud* and access it from a web browser or a designated client. Instead of storing files on physical media, it is today a common practice to use services such as Dropbox, Google Drive or Apple's iCloud for home and business matters. In November 2013, Dropbox hit 200 million users [1]. Google Drive is integrated across a large number of Google's products, as is the case with Apple's iCloud, which stores and syncs personal preferences across devices and applications.

The existing mainstream services mentioned above are centralized, which means the files and resources stored with them are placed on central servers somewhere on the Internet. This report describes the results of developing a decentralized file sharing system with focus on security and open web standards.

## 1.1 Background

Most of the existing technologies and protocols that constitute the internet, as well as the services running on it, are by design decentralized and promote the design of distributed systems. On the application layer, transfers of the hosted large files in peer-to-peer scenarios are to a growing extent conducted in a distributed fashion using the Bittorrent protocol, to the point where it is becoming the de facto approach for many use cases and areas. In others, the transition to distributed transfers and storage of data is still in its infancy. In technology and developer communities, *distributed* and *decentralized* have become buzzwords. The norm today is to use distributed approaches for things such as source version control, data storage, heavy computations and content delivery. However, both users, developers and businesses are moving more and more data to an arbitrary *cloud* on the internet. Companies such as Google and Dropbox provide servers for storage of data: A practice that poses several security concerns. Recent news on government infiltration of these services, as well as the revelation of Microsoft defense of private investigations in users' Hotmail inboxes, raises the issue of centralized storage beyond users' control. The internet itself has always been decentralized, and resembles a conventional graph structure with nodes and paths. By centralizing information and giving up a *thin server-fat client* concept, one deviates from the fundamental idea of a decentralized network.

There is currently a clear transition of user-space applications and services from native binaries to web applications running inside a web browser. Recent initiatives such as Google Chrome Apps, Adobe Phonegap and Mozilla Firefox OS are starting to bridge the gap between *web apps* and *native apps* even more, both for mobile and desktop environments. These applications are implemented using what is casually referred to as *HTML5* or, more accurately, the open web stack – an umbrella term for technologies such as HTML, CSS and JavaScript, which are based on open standards. Web applications are becoming ever-increasingly powerful in areas of software engineering and computer science, even though many standards are still in their infancy. Browser implementation and support is still unstable in many areas, leaving much to cover. Even so, technologies for functionality that was earlier exclusively for native applications are now available for any developer to use in modern, cutting edge web browsers. Notable examples are peer-to-peer video chat, local file storage, powerful encryption methods, and real time full-duplex communication.

Following these trends, a natural consequence is a peer-to-peer distributed data-syncing protocol implemented purely on the open web stack utilizing cryptographic keys for access control. This would hopefully act as a stepping stone facilitating the development of user-friendly and convenient, yet secure and privacy-protecting distributed implementations of services such as personal file syncing, media sharing and private communication.

## 1.2 Purpose

Development of a distributed peer-to-peer encrypted system for storage and communication of resources packaged in a library, with a web browser as the only client-side dependency has been conducted. In this report, the results are presented along with the state of modern web standards as these technologies are researched and analyzed in terms of how they can facilitate realization of such as system.

**Rymd** is the main outcome and end goal of the project. It will solve authentication and data transfer between nodes, while it also provides encryption and storage of resources locally. It should be usable as a drop-in module by any web client side code, such as a regular front-end web application, browser extension, or widget.

**Shuttle** is a proof-of-concept prototype using Rymd to show its functionality. It can be seen as an executable evaluation of Rymd and will be briefly discussed in this report.

Below are the main goals and requirements of Rymd:

**Privacy** . Only users that are given explicit access to a resource should be able to deduce anything useful about its content. No central entity, such as a server administrator or network operator, should be able to extract incriminating information about a client. Users should be able to trust that they know who they are communicating with. No network operator or server administrator should be able to forge identities in a way that can not be detected by a user.

**Security** . Encryption in all layers, from resource storage to data transfer.

**Reliability** . If any server goes down and can not be recovered, no damage should be done to the network as a whole as long as anyone can host a new server using the same source code.

**Modularization and agnosticism** . The system as a whole should not depend on particular implementations for resource storage, key storage or which protocol to use for data transfer. If a developer wants to, they should be able to easily plug in their own alternative implementation module.

Shuttle should be a working example of a file-sharing application that leverages Rymd to provide all these features.

## 1.3 Problem

There are several sub-problems in a system of this nature that the project needs to address.

- Decentralization of system logic
- Source code integrity
- Peer identity verification
- Resource storage
- Resource identification
- Communication flow and transfer initiation

Below, these main points affecting the architecture are presented.

### 1.3.1 Decentralization of system logic

In a truly distributed system, it is necessary to avoid having crucial system logic and data on a server. The functionality of the system should not rely on availability any central server. Temporary downtime can be accepted as long as servers do not store crucial data and can be replaced with new servers running the same software. Since clients can not know what application their peers are running, all information from other peers must be verified and considered as untrusted.

### 1.3.2 Source code integrity

Since the system logic is run in a web browser, how can one trust the fact that the client code is not altered between executions? This issue is one of the reasons why a large part of the online community is considering client-side JavaScript encryption to be a generally bad practice. At the time of this writing, major web browsers have no way to verify client-side code or resources the way that native binaries or Java applets can be cryptographically signature checked.

### 1.3.3 Peer identity verification

Each user of the system will be associated with a self-generated private-public pair of cryptographic RSA keys. With knowledge of the public keys of their peers, there are standardized identity verification protocols used on a session-to-session basis. Regardless of the authentication protocol used, there is always a chicken-and-egg problem with the distribution of public keys and how to tie them to identities. In order to trust the validity of the key provided from another entity, the user puts trust in that entity. Traditionally, there are two types of Public Key Infrastructures (PKIs) with different ways to address this:

- A Web of Trust, as often utilized in OpenPGP [2]. Here, a user has a list of peers that they trust - trusted introducers. If they receive a public key and associated identity signed by one of their trusted introducers, they will know that the trusted introducer has verified the connection between the identity and the public key. This way, an active user will steadily grow their network of trusted introducers. One needs to have a network of dependable and active peers in order to successfully participate in a Web of Trust.
- A PKI centered around one or several Certificate Authorities (CAs). Here, there is a predefined list of authorities that are trusted to sign participants public keys. This creates a centralized network and puts a lot of trust in the CAs. SSL utilizes this approach and there are several historical examples of when this trust has been broken (more recently in the Diginotar hack of 2011).

### 1.3.4 Resource storage

Usability, security and adherence to public web standards are three priorities that make the question of how to locally store resources on clients a difficult one. There are proprietary technologies for mapping resources to files on the local filesystem, which could be very useful – but without cross-platform support, they are considered out of question.

### 1.3.5 Resource identification

It is desirable for resources to have identifiers that are both memorable, secure, and unique - Zooko's triangle is a concern here as well. However, there are practical limitations on using any current cryptocurrency blockchain here. There is a monetary cost associated with the insertion of a new value, and updates can take a significant amount of time to propagate over the network. These practical issues would make such a system practically unusable. File names are not even close to unique and disclose unnecessary information, should an adversary without the corresponding secret key get hold of an encrypted resource. Since

resources are communicated peer-to-peer, the issue of malicious resource identifier collision attacks becomes negligible since users would have first-hand contact with peers that they trust and can verify the identity of. Resource checksums will have to be communicated and verified by peers before accepting a transfer of resource data.

### 1.3.6 Communication flow and transfer initiation

Once peers have verified each other's identities, the question of how two peers start the transfer of a shared resource arises. This is also an implementation detail in applications using the library and depends on what problems that particular application is intending to solve.

## 1.4 Scope

The project is scoped to encompass two parts: *Rymd* and *Shuttle* (see 1.2)

The system will not deal with version management, syncing, merging resources, and history. Neither will the upload of the users' keys (used for encrypting data) to the blockchain be solved by Rymd, since it involves payment logistics – a subsystem not in our scope.

Leaking of certain kinds of metadata will not be addressed. Namely, information on who is communicating with whom, since this is a very difficult issue far beyond the scope of this project. Also, network operators will likely be able to make a rough estimate on resource size based on the amount of data transferred, since this is considered a reasonable privacy-performance tradeoff as long as transfers are padded enough that the estimation can only ever be so vague.

Some of the technologies involved in this project are quite recently developed, which means that some of even the latest browsers might lack support. The final product will most likely not work on all types of devices and browsers.

## 1.5 Similar technologies

There is a plethora of disparate technologies for distributing and syncing data between peers that at first glance may look very similar to Rymd. Below are some more well-known and similar pieces of software. The features described will hopefully highlight how they relate to each other and Rymd.

**Bittorrent Sync** [3] is a distributed peer-to-peer multi-way file syncing software using the Bittorrent protocol for file transfers. Synced folders are mapped directly to the underlying file system, and each folder is encrypted using a shared secret key. Public-key cryptography is not employed, and there are only closed-source binary clients using Bittorrent, Inc's network available. While they do have a developer API, it requires developer keys issued from Bittorrent Inc.

**RetroShare** [4] markets itself as a Friend-2-Friend decentralized communication platform. It uses GPG to create a Web of Trust between peers. It is, however, a very large project: The application provides file-sharing, instant messaging, discussion forums, e-mail, VoIP and group chat. It is open source and distributed as cross-platform binaries.

**ShareFest** [5] is a peer-to-peer one-to-many file-sharing web based software using WebRTC data channels. ShareFest can be seen as a more limited and primitive version of what Rymd aims to be: ShareFest can share files over WebRTC channels, but does not accommodate for authentication, persistence or local encryption. It does, however, operate on a mesh network similar to Bittorrent. Other similar WebRTC-based P2P file sharing web applications but without additional cryptographic properties include RTCCopy and ShareDrop.

**Freenet** [6] is one of the first *darknets*, consisting of a distributed, decentralized data store that uploads files with strong anonymity across a network. Each node in the network also acts as a cache for the content stored in the network. Files are generally split up in parts that are distributed, and when fetching files it is unfeasible to determine the origin and sender of the files. Focusing on anonymity, free speech and plausible deniability, the encryption is done in the communication and storage layers. Because of this design, Freenet is quite slow. Files can be retrieved using the cryptographic key used to upload it. Freenet is free software built with Java.

**Tahoe-LAFS** [7], or Tahoe Least-Authority Filesystem, is a distributed, encrypted and redundant file system. It distributes encrypted files across a predetermined set of servers and allows sharing of both mutable and immutable files. There is a web-interface, but equal to all other user-interfaces it has to go through a *gateway* where encryption and server-communication is performed. Users will typically run their own gateways and will thus need to accommodate for hosting for them.

**Bitmessage** [8] is a P2P distributed messaging system intended to replace e-mail. Public keys of all participants are distributed over the entire network, and can be retrieved using their fingerprints (which are used as addresses). In order to send a message, it is encrypted using the receiver's public key and sent to the entire network. Participants try to decrypt every message, and will so be able to retrieve the ones they can decrypt. Messages are stored in the network for two days. There is thus no way to tie messages to senders and recipients.

## 2 | Methodology

Initially the project was split into two parts: an evaluation phase and an implementation phase. This chapter intends to further detail the results of each phase and how it relates to the end product.

### 2.1 Evaluation of technologies

The goal of the evaluation phase was to map out the landscape of relevant technologies. Different options were compared to each other in order to analyze strengths and weaknesses in regards to a set of given parameters:

**Suitability** . How well does the technology suit the needs and demands of the job? Are there any technical limitations?

**Maintenance** . Is the technology actively maintained? If not, does it pose an issue? What are the future scenarios?

**Industry support** . Are some unsupported browsers negligible? What are the industry's current opinions?

Research was made about open web technologies, mainly belonging to the HTML5 standard. The usage of open web technologies was a requirement of the project's end product, Rymd, which meant that no native code could be written as part of the system and that the quality of the product would be completely dependent on the state of existing APIs and tools of web development. Thus research was made in order to survey the landscape of existing technologies at the time in order to determine which, if any, fulfilled the requirements so that the end product actually could be developed using them.

A set of areas were created, where each area connected with one or several core problems stated within the project. In each area, evaluation and comparisons were made, which included researching APIs and prototyping actual test cases implementing isolated forms of future system features. The research areas were divided as follows:

**Data storage** . Investigated how to store data locally for each node. This was crucial in order to meet with the overall requirement of building a decentralized system.

**Communication** . Reviewed the possibilities for communicating and sending data with peer-to-peer technology between two nodes.

**Authentication and Permissions** . How to solve authentication for nodes.

**Prototyping** . Rapidly produced a rough test case for sending a file from one node to another.

## 2.2 Prototyping

The aim for the prototyping area was to quickly decide if it was in any way possible to achieve the requirements with the technologies chosen. Thus was a rough prototype of Rynd and Shuttle created, which implemented two basic test cases: storing a file in the chosen data storage implementation, and sending that file to another node where it was stored in that node's local data storage. The prototype worked successfully, which validated the choice of those special technologies.

## 2.3 Implementation

During the implementation phase, which stretched from the end of the evaluation to the end of the project, the product was implemented according to the requirements. It was early decided that the implementation process would lightly apply agile methodologies. For this project, this included having a Product Backlog with User Stories, working in sprints, and having bi-weekly Scrum-meetings where current state and eventual problems were brought up. For managing the stories in the backlog, the online management tool Pivotal Tracker was used.

All source code was managed by the distributed source versioning system git<sup>1</sup> and with the online tool GitHub<sup>2</sup>. It was decided to split up the Rynd library into several smaller code repositories ("modules"), which all resided on GitHub. See section 4 for details regarding the technical implementation.

# 3 | System design

This chapter describes the technical foundations of the system and how the problems described in section 1.3 have been addressed in Rynd.

## 3.1 Overall system design

### 3.1.1 Resources

## 3.2 Peer identity verification

For a truly decentralized system, it is not acceptable to adapt a CA-entered approach. While a Web of Trust is interesting, it might be too cumbersome for users. This issue is addressed in "Zooko's Triangle" (See figure 3.1), stating that no system assigning names to participants in a network can have the property that names are secure, decentralized and meaningful at the same time. This conjecture has since been proven false by the design of systems such as the blockchain of the cryptocurrency Namecoin, which effectively acts as a cryptographically

---

<sup>1</sup><http://git-scm.com/>

<sup>2</sup><https://github.com/rymdjs>



secured distributed hash table (DHT) with unique keys. Users can reserve a name and assign to it a value of their choice by the cost of a small amount of the Namecoin currency (at the time of this writing 0.01 NMC [9], which equals roughly 0.03 USD [10]).

Rymd therefore utilizes a DHT for storage of keys to achieve all of these goals: The distributed nature of cryptocurrencies makes it decentralized; peers can choose their own names (identities), giving meaningful names; the small monetary fee required to register a name makes it both secure and prevents massive name-squatting by malicious third parties.




Figure 3.1: *Zooko's Triangle*, with the edges representing the achievable combinations of features [11]

### 3.3 Decentralization

Assuming that the system can utilize a DHT such as a cryptocurrency blockchain for storage of the public part of RSA key pairs, the issue of how to interface a web application with the blockchain in a way that allows for verification of identities without putting too much trust in the HTTP/cryptocurrency gateway also needs to be addressed. Additionally, as previously stated, the initial insertion of the key requires monetary resources, and is perhaps something that should be solved outside of Rymd. While the public key can be stored in a DHT, private keys need to be stored securely on each client, preferably without giving client code any direct access to the raw keys. How the initial generation of keys are performed also needs consideration. Finally, a secure way to store the encryption keys for encrypted resources needs to be addressed. The question of how these resource-associated secret keys are distributed is deemed an implementation-specific question and will be out of scope for Rymd, but handled in Shuttle.

### 3.4 Modularity

The system was designed from the ground up to use interchangeable parts. This was achieved by identifying key features and separating them into individual modules. These were then used and intertwined in a central hub.

Advantages

Dependency injection

## 4 | System implementation

Standards and technologies for web applications are being rapidly developed, and the boundaries for what is possible to achieve in a web application is being continuously pushed by browser vendors and standards groups. Of relevance to Rymd, the Web Real-Time Communications (WebRTC) protocol [12] has become usable for arbitrary data streams in major browsers in 2014. Also of relevance, the still unfinalized Web Crypto API [13] is currently available in an experimental stage in recent months at the time of this writing. The goals of Rymd has thus become technically viable in a web environment as of very recently.

Furthermore, there is currently a lot of work underway in the field of distributed secret communication. Notable projects that share similar ideas or have inspired Rymd are mentioned under section 1.5. Also worth noting is the new field of cryptocurrencies, which work by distributing a cryptographically based ledger over an entire network. The first and most well-known is Bitcoin [14], but there are also subsequent currencies that extend the original idea beyond that of a traditional currency to a system that can be used for a wide range of applications. The first worth noting is Namecoin [15], which adds a global and secure key-value store. Another interesting initiative is Ethereum [16], a cryptocurrency with contracts that not only allow storage of arbitrary data in the blockchain, but can also be scripted with a Turing complete programming language and can therefore be used to implement arbitrary systems. A system like Ethereum could be very interesting to explore for a project like Rymd, but it is still in such an early stage that it is deemed too unstable to be useful at this point. Namecoin is currently considered a good candidate for key distribution.

Keybase [17] is another recent initiative that intends to solve the distribution of public keys. It is essentially a HTTP-interface that maps keys to identities. While commandable, it again raises the issue of centralized storage. Systems relying on Keybase put a lot of trust on the availability and integrity of their service.

### 4.1 Data storage

Storage of data is, suffice to say, crucial for any file sharing system. Since the data store was to be used by several parts of the application the demands for the module's interface had to be as general as possible, adhering to a standard CRUD<sup>1</sup> interface, including methods for creating, fetching, updating and deleting records in the store.

There are essentially four alternatives for persisting data on the client:

- LocalStorage
- IndexedDB
- WebSQL
- FileSystem API

LocalStorage is included in the HTML5 WebStorage specification [18], and is a basic key-value store with a simplistic API. It is supported across all major browsers and has a maximum storage limit of 5 megabytes. The latter was a deal-breaker, since the product would have to support larger files than could possibly fit into that space. LocalStorage further does not support complex structures and indexing, and storing different data types is complex and needs manual serialization and deserialization. Thus this solution was rejected at once.

IndexedDB and WebSQL are both client side databases and more sophisticated storage solutions than LocalStorage. WebSQL is supported by Google Chrome, Apple Safari (desktop and iOS), Opera and Android. The specification is not longer maintained by W3C [19], and will probably be deprecated on all browsers in the future. IndexedDB is supported by all

---

<sup>1</sup>Create–Read–Update–Delete

	IndexedDB	WebSQL	File System	LocalStorage
Google Chrome	Yes	Yes	Yes	Yes
Mozilla Firefox	Yes	No	No	Yes
Apple Safari	No	Yes	No	Yes
Opera	Yes	Yes	Yes	Yes
Microsoft Internet Explorer	Yes	No	No	Yes

Table 4.1: Browser support for selected HTML5 APIs

major browsers except for Safari (desktop and iOS), and is a Candidate Recommendation by W3C [20]. Arbitrary types of data can be stored in the database, such as strings, numbers, Javascript objects, and raw binary data.

The FileSystem API is a collection of methods for reading and writing to a sandboxed filesystem from a browser with client Javascript code. It is a very early standard, and is currently only supported by Google Chrome and Opera, and has the status of Working Draft by W3C [21]. While FileSystem has good performance for larger files and a well-performing asynchronous API, it lacks support for indexing and search. Mozilla seems to have no plans on implementing FileSystem in Firefox [22].

All of the mentioned technologies are sandboxed: the data is tied to a single origin (*http://test.domain.com* for instance). All future access to the data must come from that domain (this includes the protocol and port number as well). The browsers also limits the maximum allowed storage size – the quota. The quota is different for each storage mechanism, and the browser typically asks the user with a dialog if they want to let the app exceed the quota.

The conclusion was to use IndexedDB for persisted resource storage. It was chosen because of its support by Google Chrome, Internet Explorer and Firefox, and due to the fact that its actively maintained (while WebSQL is not). Users of the Safari browser will not be able to utilize the product, but in respect to the project’s overall direction in regards to experimental technologies, this is negligible.

#### 4.1.1 IndexedDB

The IndexedDB is a transactional, indexed client side database capable of storing different types of data structures with an asynchronous API. IndexedDB is actively developed and implemented in the latest versions of Mozilla Firefox, Google Chrome, Microsoft Internet Explorer and Opera: its specification is a Candidate Recommendation by the W3C, as of July 2013[20]:

This document defines APIs for a database of records holding simple values and hierarchical objects. Each record consists of a key and some value. Moreover, the database maintains indexes over records it stores. An application developer directly uses an API to locate records either by their key or by using an index. A query language can be layered on this API. An indexed database can be implemented using a persistent B-tree data structure.

#### Basic structure

Due to IndexedDB’s object-oriented nature, a database includes a set of *object stores*, which act similar to tables in relational database management systems. An object store can hold *objects* of different types, including binary data and Javascript primitives and objects. Each object has a *key* (either specified by the developer from the object’s properties, or automatically generated and managed by the database) that is used for indexing and retrieving records. One or several *indexes* can be created on a store from an object’s properties for quick querying. A *cursor* is used to iterate on the resulting set of objects from a query on the store.

The asynchronous API might include complex patterns if the developer is not used to NoSQL structures. Unlike WebSQL, IndexedDB does not support SQL, and instead exposes

ways for querying and manipulating data via *requests* and *transactions* (see section 4.1.1). A positive facet of the rejection of SQL in favor of NoSQL is the prevention of SQL injection attacks, but with the cost of a steeper learning curve for already experienced database developers. Queries to the database will not yield the resulting data set: instead requests are returned, which will trigger *events* for when the operation is finished. When an event is triggered a *callback* can be passed to handle the scenario and use the data. This goes well with the asynchronous nature of Javascript, where events and callbacks are used heavily. Using asynchronous passing of callbacks prevents the program execution to block at a line when a potentially heavy operation is called. The Javascript code snippets below show the difference in synchronous and asynchronous calls.

```
// Fetch a record with id 10 from a database and store in variable
var result = DB.find(10);
```

Listing 4.1: Synchronous call

```
/*
  Request a record with id 10 from a database, continue execute other code,
  and handle result of the database operation in callback when it has finished
*/
var request = DB.find(10);

request.onsuccess = function(evt) {
  var result = evt.result;
};
```

Listing 4.2: Asynchronous call

## Security and reliability

IndexedDB is built on a transactional model, which implicates that all commands runs inside a transaction context. Transactions have a certain lifetime, and cannot be used after its expiration. This transactional model is especially useful for when several instances of a client application is using the same database and issuing commands: without transactions, concurrency problems and other collisions might occur with data loss as a result. Transactions are able to abort and be rolled back to the state of the database before the transaction was started.

Kimak, Ellman and Laing highlight four important aspects of securing a IndexedDB driven application in their *An Investigation into Possible Attacks on HTML5 IndexedDB and their Prevention*[23]:

- Client side data encryption
- Input validation
- SOP (Same-Origin Policy)
- Code analysis

The database in IndexedDB does not include any kind of bundled encryption or validation, which means it is the developer's responsibility to sanitize and encrypt sensitive data before inserting into the store. Encryption is vital for the scenario where the contents of the database is compromised: the attacker must have access to the encryption key in order to read the information in plain text. Validation is needed if malicious content, such as Javascript, is inserted as the data fields in the store and then will be executed at a later stage.

The *Same Origin Policy* is used in IndexedDB. An origin is the transfer protocol, the domain, and the port number. Thus every database is associated with an origin, which implicates certain security aspects: an application in `http://domain.com/subdir` may retrieve data from `http://domain.com/subdir/dir` since they have the same origin, but cannot retrieve data from `https://domain.com:3000` due to the different protocol and port number. This is a layer of protection against Cross Site Scripting (XSS) attacks, even though there is no prevention against XSS holes in the other parts of the application (the database might be compromised due to malicious scripts injected elsewhere).

Code analysis is divided into *static* and *dynamic* analysis. Static analysis is the analysis of the to-be insterted data in order to detect malicious material. Dynamic analysis is the analysis of executed programs, and can according to [23] be done by checking the call from the web application to the database and on success, the database operation can be performed.

### 4.1.2 Implementation

The main task for the data storage module was to abstract away the low-level methods in IndexedDB (the backing store used, see section 4.1.1). By the use of the concept of *promises*<sup>2</sup>, the asynchronous, callback-based methods in the IndexedDB API could be made very streamlined and simple to manage.

```
var Store = new IndexedDbStore("myStore")

// Fetch all records as an array
Store.all().then(function(records) { ... })

// Create a record
Store.create("A record").then(function(record){ ... })

// Insert a record
Store.save("A record").then(function(guid){ ... })

// Fetch a record by GUID
Store.get(guid).then(function(record) { ... })

// Delete a record by GUID
Store.destroy(guid).then(function(record) { ... })
```

Listing 4.3: Common database operations

The largest challenge came to the edge cases when storing files, or as they are called in web browser: *Blobs*. Since only Firefox can as of now store blobs directly in IndexedDB, an alternate route had to be taken for other browsers. Initially the module used conditionals and converted incoming data to and from *ArrayBuffers* (the browser construct for raw byte streams). But since *ArrayBuffers* are just the raw data, all metadata for the blobs (such as filename, timestamps, size) would be lost when saving as an *ArrayBuffer*. In early versions of the data storage module this metadata would be stored in a separate store in the database, but this was too tight coupled and was removed. The final implementation is storing data as-is – any metadata must be saved explicitly in a separate operation.

The data storage module is a separate repository and the source is available at <https://github.com/rymdjs/data-storage>.

---

<sup>2</sup>[http://en.wikipedia.org/wiki/Promise\\_\(programming\)](http://en.wikipedia.org/wiki/Promise_(programming))

## 4.2 Communication

In order for nodes to be able to share data, they need a way to connect to each other. They also need to do this in a secure manner in order to prevent potential vicious third parties listening on a connection from making any sense of retrieved data. I.e. critical parts should not be sent in raw form but rather be encrypted. When considering security aspects there are essentially three questions that need to be answered regarding the issue of connecting nodes:

- How can a node find another node to begin with (peer discovery)?
- When a node has been found, how can a connection be established?
- What data needs to be encrypted in order to ensure the integrity of the system?

Answering these questions and finding the corresponding best technology was the focus of this research area. There are a bunch of different emerging trends on the web today and some of them enable peer-to-peer communication.

WebRTC seeks to be a common standard for browsers with W3C drafting client side APIs and IETF developing the protocols and peer-to-peer communication[24]. There are also security aspects built into WebRTC. The major browser vendors Google, Mozilla and Opera support the project [25]. While Microsoft actually supports the concept of WebRTC and contributes to the W3C WebRTC working group, it does not support Google's (or nowadays, W3C's and IETF) version of it[25].

Microsoft warns about supporting the new technology until it has actually become a standard and also doesn't fully agree on some constraints put on the technology[25]. Microsoft explains that one of their issues with the current WebRTC version is that it has predetermined paths of choosing codecs and ways of sending media over the network – sort of similar to a black box. This hinders application developers wanting to optimize to suit their own needs. Microsoft's answer to this is their own CU-RTC-WEB (Customizable, Ubiquitous Real Time Communication over the Web) which tries to address these issues.

All in all Microsoft remains optimistic that a common standard will eventually be established[25]. They do however stress that more participants need to get involved besides Google and Mozilla. Since all major browser vendors (except Microsoft) support WebRTC at this time we chose this technology for the project.

### 4.2.1 WebRTC

WebRTC is a project which enables real-time communication between browsers[12]. Through the project, developers are able to create different types of applications which leverage peer-to-peer technology.

To obtain and transfer streaming data WebRTC's functionality is abstracted into three different APIs: *MediaStream*, *RTCPeerConnection* and *RTCDataChannel*[26]. *MediaStream*, or *getUserMedia*, consists of synchronized media streams, e.g. synchronized video and sound from a computer's camera and microphone. *RTCPeerConnection* manages reliable and efficient communication of the data streams. There are a lot going on under the hood that ensures reliability for real time communication even in unstable networks. Furthermore, when a peer connection is present arbitrary data can be sent by leveraging *RTCPeerConnection* with *RTCDataChannel*.

#### **RTCDataChannel**

The *RTCDataChannel* demonstrated the capabilities that Rymd desired. Data transfers are secured with the DTLS (Datagram Transport Layer Security) protocol. The DTLS protocol is based on the TLS (Transport Layer Security) protocol, the main difference being that DTLS is constructed for datagrams while TLS is used for more reliable transport protocols such as TCP.

Before a connection can be initiated between peers, one of two parts must extend an offer which contains data describing the connection to the other part - this is often referred to as the signaling phase. The signaling phase requires a channel where the offer can be negotiated - it is usual for the channel to be a dedicated signaling server, but examples of a more serverless approach can be found[27]. The standard does not provide any recommendations regarding the choice of signaling channel and protocol - this is for developers to decide. The connection phase is then handled by `RTCPeerConnection` which manages the ICE (Interactive Connectivity Establishment) workflow.

## ICE

ICE is a framework with the objective of connecting peers[26]. Since computers today are usually behind NAT gateways they are using different IP-addresses in their private network and the same IP-address publicly to access the Internet. This poses difficulties when trying to achieve a peer-to-peer connection.




Figure 4.1: *The different ways for ICE to find network interfaces and ports [26]*

ICE decides between different ways of achieving a connection and chooses the one with the lowest latency[26]. At first ICE tries to connect peers cleanly through UDP and STUN (Session Traversal Utilities for NAT). STUN servers are simple, cheap to run and enables a peer behind a NAT to discover its public address and port. When the IP-addresses are known data doesn't need to go through the server but instead flows directly between the peers. However, if it is not possible to establish a connection ICE attempts to use TCP instead of UDP: first HTTP and then HTTPS.

If STUN doesn't work, probably because of firewalls and more advanced NAT systems, ICE uses cloud fallback in the form of a TURN (Traversal Using Relay NAT) server. This connection isn't peer-to-peer since the data needs to go through the server and because of this it is also slower than the previous alternatives. But it allows a connection to be made under almost any condition.

### 4.2.2 Implementation

Rymd leverages the open source project PeerJS, which simplifies sending peer-to-peer data between clients. PeerJS makes use of WebRTC and is essentially split into two components: A server that acts as the signaling channel and a client side API which interacts with the server as well as other peers. The server just handles the brokering of connections, which implies that only the data necessary for negotiating a connection is sent through that point.




Figure 4.2: *STUN servers lets peers in a private network behind firewalls discover their public IP-adresses[28].*

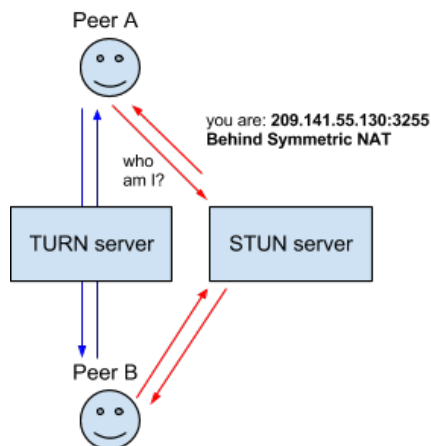


Figure 4.3: *If a peer-to-peer connection cannot be made, a relay through a TURN server could be used instead. All peers then send their packets through the relay which is more costly but makes the system work[28].*



After a connection has been setup between two clients, the server is no longer needed in order for them to communicate.

Behind the scenes Peer.js handles the communication between clients and the server with the help of Websockets, but falls back to other protocols if a browser lacks support.

## 4.3 Authentication

### 4.3.1 Namecoin

A phenomenon that has been on the rise during recent years is that of cryptocurrencies such as Bitcoin [29]. Each participant in the Bitcoin network keeps a ledger of all transactions throughout the history of the network. In order for a transaction to be deemed valid, it needs to be included in a cryptographically signed block together with a salt by one of the nodes, with a hash that has a special format. It is this brute-force search for salts that generate these hashes that are called *mining* and constitute the work done by *miners* to keep the network running. As an incentive, each verified block also includes a reward to the miner that first finds it and submits it to the blockchain[30]. In effect, all transactions ever made are publicly available and tracked so that anyone can confirm their validity. This prevents forgery and double-spending of bitcoins.

Another cryptocurrency is Namecoin[29]. Namecoin is essentially a fork of Bitcoin with new transaction types that allows its blockchain to be utilized as a distributed key-value store. Although similar in nature to Bitcoin, its main purpose is to be used as a decentralized domain name system (DNS), rather than as a monetary currency.

With a decentralized DNS such as Namecoin, top level domains (such as *.com* or *.se*) can exist without being controlled by any central authority [29]. Also, the DNS lookup tables where domain names and their IP addresses are stored are shared in a peer-to-peer manner. The only necessary condition for these domains to be accessible is that there are participants willing to run the DNS server software. Although mainly intended to be used as a DNS, it contains several namespaces where arbitrary strings such as public cryptographic keys can be stored.

## 4.4 Encryption

## 4.5 Testing

Automatic unit tests have been implemented where possible. No integration or functional tests have been written for testing larger parts of the system, due to the fast iteration of the library's interface and constant change in the implementation.

# 5 | Results

This project has resulted in a modular peer-to-peer developer library for sending data, encrypted with public-private encryption, over a secure connection to another client. The library, *Rymd*, has composed several modules for specific areas such as data storage (section 4.1), cryptography, (section ??) and communication (section 4.2.2) in order to create a foundation for sending files without central file storage.

For demonstrating the capabilities of *Rymd*, a sample prototype web application has been created, named *Shuttle*. This front facing client provides a user interface for showing local files, sending files to other users registered in the blockchain, and encryption key management. In *Shuttle*, the user is capable of adding, listing, viewing and deleting files in their local data store, where the files are encrypted and stored along with their metadata. By knowing a recipient's username, the user is able to share a file with the recipient over an encrypted P2P connection. When sharing a file, the receiving end will instantly show a notification

with a remark that the sender wants to share a file. If the recipient chooses to accept the sharing request, the file will download to their local data store.

## 5.1 Source code

Due to the system's modularity, a number of repositories exist to hold the source code of the different modules. All source code is available at the project's GitHub page: <https://github.com/rymdjs>. All relevant repositories can be found below:

**Rymd** <https://github.com/rymdjs/rymd>

**Shuttle** <https://github.com/rymdjs/prototype>

**Crypto** <https://github.com/rymdjs/crypto>. Cryptography module.

**Data Storage** <https://github.com/rymdjs/data-storage>. IndexedDB adapter.

**PeerJS Connection** <https://github.com/rymdjs/peerjs-connection>. Connection adapter for PeerJS for doing P2P.

**DHT Client** <https://github.com/rymdjs/dht-client>. Client module for Namecoin lookups.

**DHT** <https://github.com/rymdjs/dht>. HTTP REST adapter for lookups in the Namecoin blockchain.

**Rymd Logger** <https://github.com/rymdjs/rymd-logger>. Custom debug and flow logger module.

**Rymd Utils** <https://github.com/rymdjs/rymd-utils>. Globally used util functions.

## 6 | Discussion

### 6.1 Quality of the resulting product

## 7 | Conclusion

# References

- [1] J. Constine. (Nov. 2013). Dropbox Hits 200M Users, Unveils New "For Business" Client Combining Work And Personal Files, [Online]. Available: <http://techcrunch.com/2013/11/13/dropbox-hits-200-million-users-and-announces-new-products-for-businesses/>.
- [2] U. Maurer, "Modelling a public-key infrastructure", *Computer Security — ESORICS 96*, ser. Lecture Notes in Computer Science, E. Bertino, H. Kurth, G. Martella, and E. Montolivo, Eds., vol. 1146, Springer Berlin Heidelberg, 1996, pp. 325–350, ISBN: 978-3-540-61770-9. DOI: 10.1007/3-540-61770-1\_45. [Online]. Available: [http://dx.doi.org/10.1007/3-540-61770-1\\_45](http://dx.doi.org/10.1007/3-540-61770-1_45).
- [3] BitTorrent Inc. (Feb. 2014). BitTorrent sync, [Online]. Available: <http://getsync.com/>.
- [4] RetroShare Team. (Feb. 2014). RetroShare, [Online]. Available: <http://retroshare.sourceforge.net/team.html>.
- [5] Peer5. (Feb. 2014). Sharefest, [Online]. Available: <https://www.sharefest.me/>.
- [6] I. Clarke. (Feb. 2014). Freenet, [Online]. Available: <https://freenetproject.org/>.
- [7] Tahoe-LAFS. (Feb. 2014). Tahoe-LAFS, [Online]. Available: <https://tahoe-lafs.org/trac/tahoe-lafs>.
- [8] Bitmessage Community. (Feb. 2014). Bitmessage, [Online]. Available: [https://bitmessage.org/wiki/Main\\_Page](https://bitmessage.org/wiki/Main_Page).
- [9] Namecoin wiki. (Feb. 2014). Register and configure .bit domains, [Online]. Available: [https://wiki.namecoin.info/index.php?title=Register\\_and\\_Configure\\_.bit\\_Domains&oldid=36](https://wiki.namecoin.info/index.php?title=Register_and_Configure_.bit_Domains&oldid=36).
- [10] CryptoCoin Charts. (May 2014). Nmc/usd - namecoin / us dollar today charts and orderbook from kraken, [Online]. Available: <http://www.cryptocoincharts.info/v2/pair/nmc/usd/kraken/today>.
- [11] Z. Wilcox-O'Hearn. (Oct. 2001). Names: Distributed, Secure, Human-Readable: Choose Two, [Online]. Available: <http://web.archive.org/web/20011020191610/http://zooko.com/distnames.html>.
- [12] Paul Gil. (2014). WebRTC, [Online]. Available: <http://www.webrtc.org/>.
- [13] M. W. Ryan Sleevi. (2014). W3C Web Cryptography draft, [Online]. Available: <http://www.w3.org/TR/WebCryptoAPI/>.
- [14] Bitcoin.org. (2014). Bitcoin – Open source P2P money, [Online]. Available: <https://bitcoin.org/en/>.
- [15] Namecoin.info. (2014). Namecoin, [Online]. Available: <http://namecoin.info/>.
- [16] Ethereum.org. (2014). Ethereum, [Online]. Available: <https://www.ethereum.org/>.
- [17] Keybase.io. (2014). Keybase, [Online]. Available: <https://keybase.io/>.
- [18] I. Hickson. (2014). W3C Web Storage draft, [Online]. Available: <http://www.w3.org/TR/webstorage/>.
- [19] —, (2014). W3C Web SQL Database draft, [Online]. Available: <http://www.w3.org/TR/webdatabase/>.
- [20] E. G. A. P. J. O. J. B. Nikunj Mehta Jonas Sicking. (2014). W3C Indexed Database draft, [Online]. Available: <http://www.w3.org/TR/IndexedDB/>.
- [21] E. Uhrhane. (2014). W3C File API: Directories and System, [Online]. Available: <http://www.w3.org/TR/file-system-api/>.
- [22] J. Sicking. (2014). Why no FileSystem API in Firefox?, [Online]. Available: <https://hacks.mozilla.org/2012/07/why-no-filesystem-api-in-firefox/>.
- [23] D. C. L. Stefan Kimak Dr. Jeremy Ellman. (2014). An Investigation into Possible Attacks on HTML5 IndexedDB and their Prevention, [Online]. Available: <http://www.cms.livjm.ac.uk/pgnet2012/Proceedings/Papers/1569607913.pdf>.
- [24] H. A. François Daoust Dominique Hazaël-Massieux. (2013). Web Real-Time Communications Working Group Charter, [Online]. Available: <http://www.w3.org/2011/04/webrtc-charter.html>.

- [25] J. Roettgers. (2012). Microsoft commits to WebRTC – just not Google’s version, [Online]. Available: <http://gigaom.com/2012/08/06/microsoft-webrtc-w3c/>.
- [26] S. Dutton. (2012). Getting Started with WebRTC, [Online]. Available: <http://www.html5rocks.com/en/tutorials/webrtc/basics/>.
- [27] Chris Ball. (May 2013). WebRTC without a signaling server, [Online]. Available: <http://blog.printf.net/articles/2013/05/17/webrtc-without-a-signaling-server/>.
- [28] C. M. Laurent Jouanneau Jérémie Patonnier. (2014). Introduction to WebRTC architecture, [Online]. Available: [https://developer.mozilla.org/en-US/docs/Web/Guide/API/WebRTC/WebRTC\\_architecture](https://developer.mozilla.org/en-US/docs/Web/Guide/API/WebRTC/WebRTC_architecture).
- [29] Crypto Coin Insider. (May 2014). Namecoin, [Online]. Available: <http://www.cryptocoinsinsider.com/namecoins/>.
- [30] Paul Gil. (2014). What Are Bitcoins? How Do Bitcoins Work?, [Online]. Available: <http://netforbeginners.about.com/od/b/fl/What-Are-Bitcoins-How-Do-Bitcoins-Work.htm>.
- [31] David Gilson. (2013). What are Namecoins and .bit domains?, [Online]. Available: <http://www.coindesk.com/what-are-namecoins-and-bit-domains/>.

# Appendices

## A | Communication flow




Figure A.1: Alice and Bob have added each other's identities as friends. Alice has a resource  $R_1$  that she wants to share with Bob.  $R_1$  is encrypted with a symmetric secret key.




Figure A.2: Alice and Bob both independently contact an HTTP gateway to get each other's public keys from the DHT. If they do not want to trust a third party, they could set up their own gateway or even add the keys manually.

## B | List of external libraries

Listed below are external tools and libraries that were used in the development of Rymd and Shuttle.

**Backbone** <http://backbonejs.org>. Lightweight Javascript frontend framework.

**Browserify** <http://browserify.org>. Node style dependency management in the browser.

**Chai** <http://chaijs.com>. Assertion library used with the test framework.

**gulp** <http://gulpjs.com>. Build system.

**jQuery** <http://jquery.com>. Javascript library for DOM manipulation and more.




Figure A.3: Alice and Bob tells a central server that they want to communicate, for example by sending their friend lists. The server sets up a direct WebRTC connection between them. They use this channel to authenticate using standard public key cryptography.

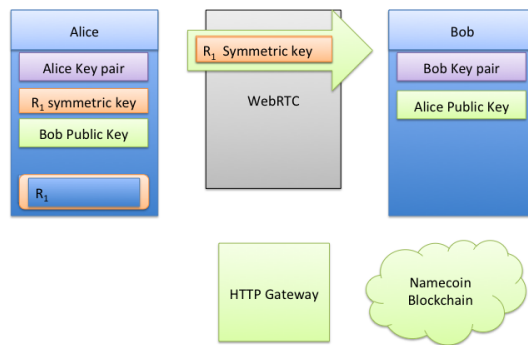


Figure A.4: Alice wants to share  $R_1$ , so she encrypts the key used to encrypt it with Bob's public key and sends it to Bob along with a reference ID that he can use to request it from Alice.

**Mocha** <http://visionmedia.github.io/mocha>. Test framework for Javascript.

**PeerJS** <http://peerjs.com>. WebRTC library for P2P communication.

**Q** <http://documentup.com/kriskowal/q>. Library for working with promises for asynchronous code.

**Underscore** <http://underscorejs.org>. Functional helper toolchain for Javascript.

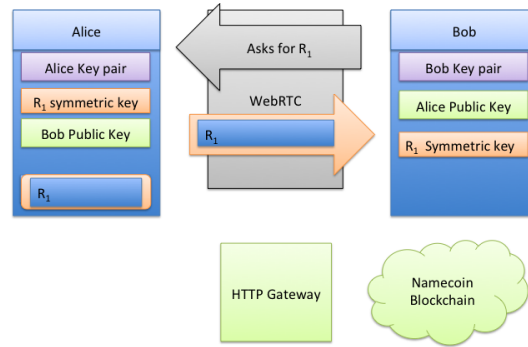


Figure A.5: *Bob asks Alice for  $R_1$ . She responds by sending it, still encrypted with the key she just sent to Bob.*

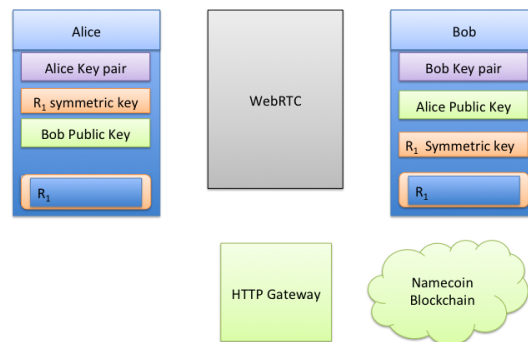


Figure A.6: *The transaction is finished. Bob is now in possession of  $R_1$  and the key used to encrypt it.*